

2025



Lesson 08

Low Code LLM App Builder (Part 2)



Tokens and Embeddings

Tokenization (Recap)

- The computers aren't very good with words but are great with numbers.
- **Tokenization** breaks down a continuous stream of text into smaller, discrete units called tokens.
- Token helps to provides a structured way to break down text into manageable pieces for the model to process.
- Tokenizer is an integral part of a language model used during training.
- Tokens are basic units that the model processes and different tokenizers can produce different token sequences for the same text.
- Understanding token is important because the service providers will charge you by tokens for API calls.

GPT-4o mini

GPT-4o mini is our most cost-efficient small model that's smarter and cheaper than GPT-3.5 Turbo, and has vision capabilities. The model has 128K context and an October 2023 knowledge cutoff.

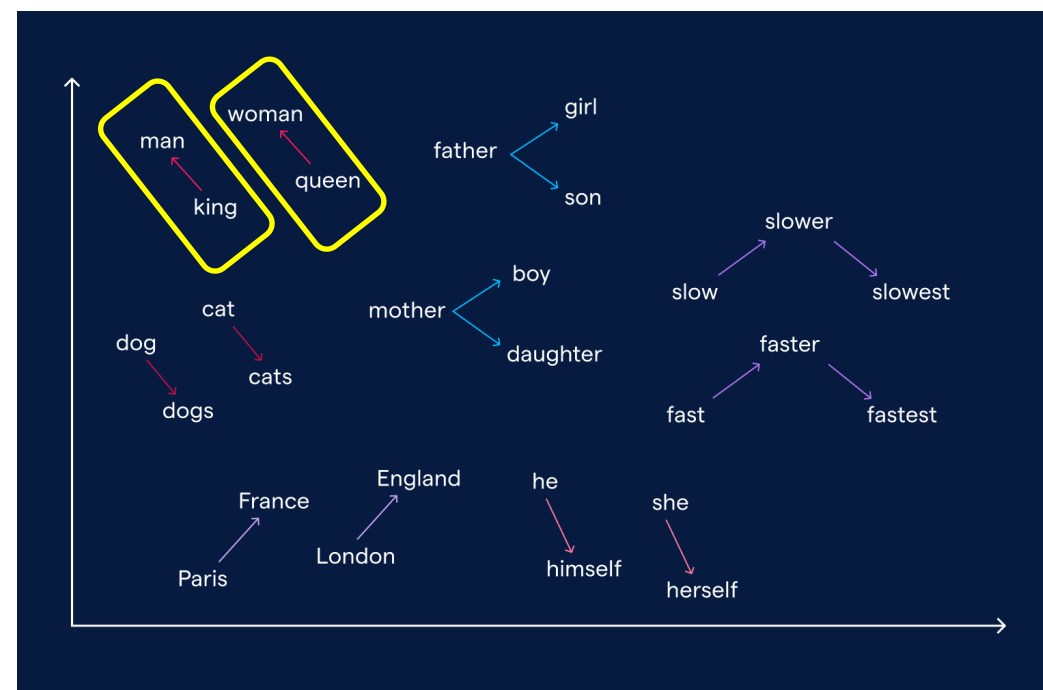
[Learn about GPT-4o mini ↗](#)

Model	Pricing	Pricing with Batch API*
gpt-4o-mini	\$0.150 / 1M input tokens	\$0.075 / 1M input tokens
	\$0.075 / 1M cached** input tokens	
	\$0.600 / 1M output tokens	\$0.300 / 1M output tokens
gpt-4o-mini-2024-07-18	\$0.150 / 1M input tokens	\$0.075 / 1M input tokens
	\$0.075 / 1M cached** input tokens	
	\$0.600 / 1M output tokens	\$0.300 / 1M output tokens

Source: <https://openai.com/api/pricing/>

Embedding (Recap)

- The next step after tokenizing the text is to convert the words to numbers.
- Transform text into something that machines can interpret is called “word embedding”.
- Word embeddings are numerical illustrations of a text.
- Sentence and texts include organized sequences of information.
- Goal is to develop a representation of words that capture their meanings, meaningful connections plus other sorts of situations in which the words are employed.

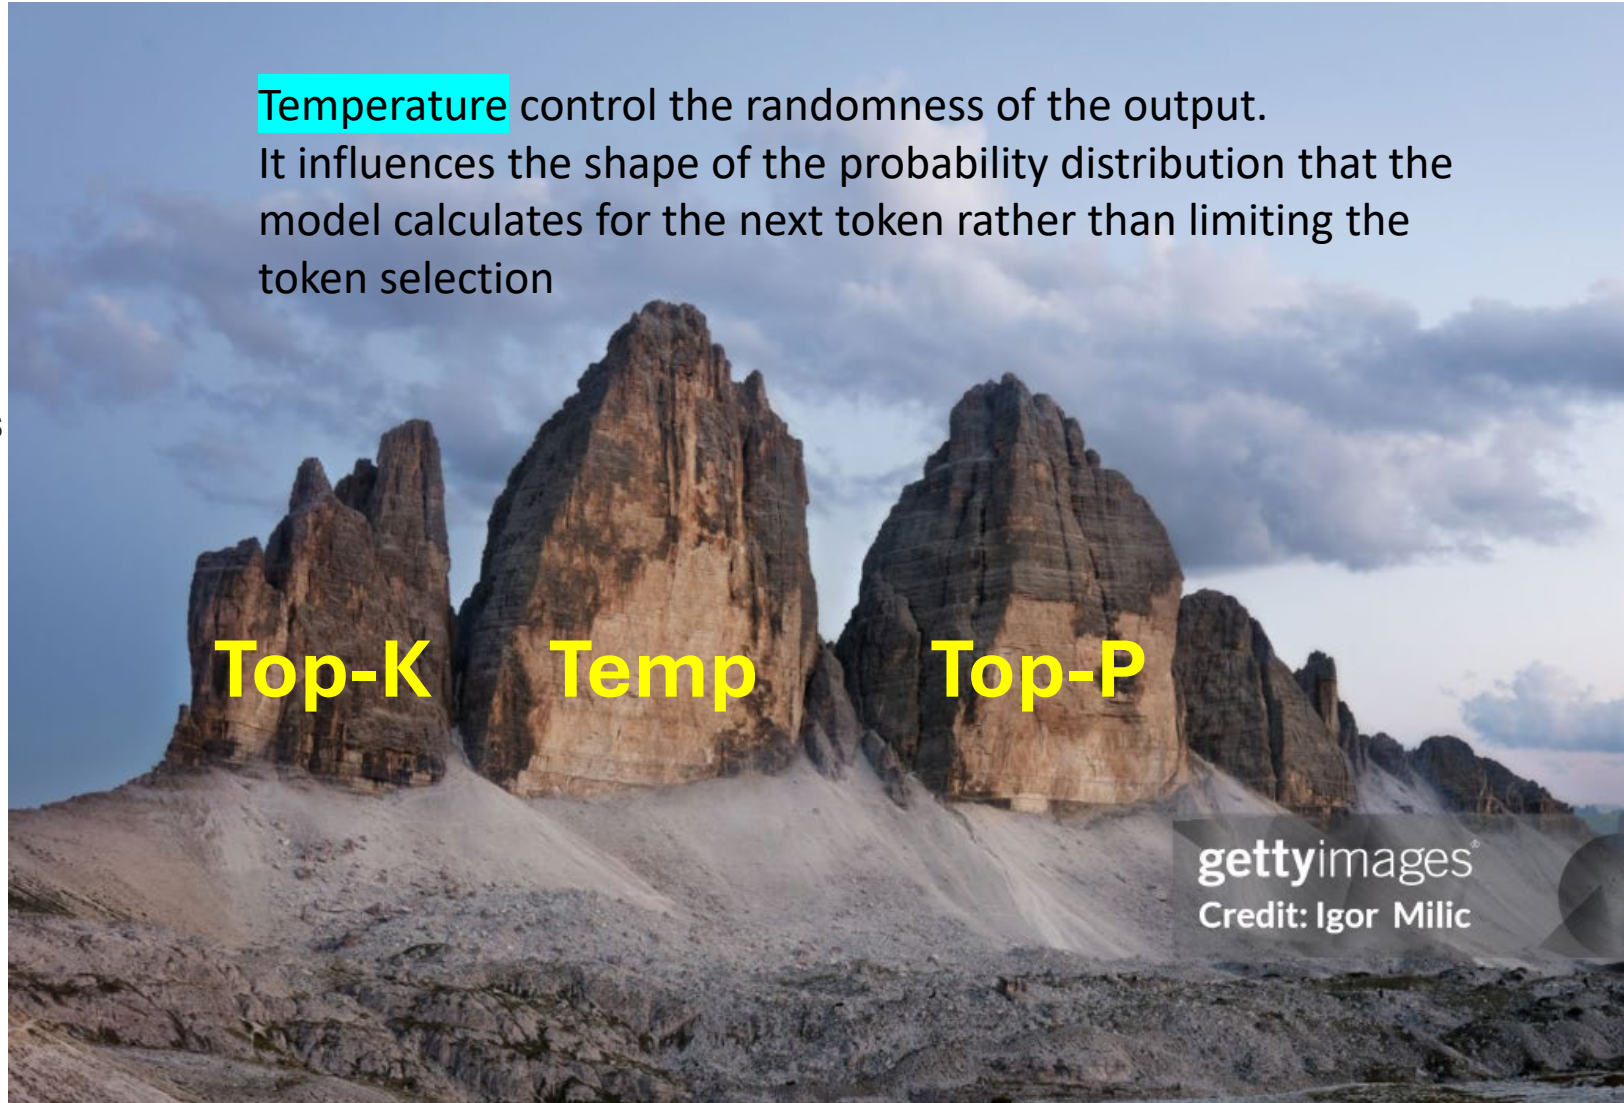


LLM Randomness

Top-K provides a controlled randomness by considering a fixed number of top probable tokens.

Temperature control the randomness of the output. It influences the shape of the probability distribution that the model calculates for the next token rather than limiting the token selection

Top-P allows for dynamic control of the number of tokens considered, leading to different levels of diversity in the generated text.



Top-K

- Limit the model's predictions to the top k most probable tokens at each step of generation.
- Instructing the model to consider only the **k most likely tokens**, truncating the less likely ones.
- For example, if k=5, the model predicts the probabilities [0.1, 0.3, 0.2, 0.15, 0.05, 0.2], it will only consider the top 5 probabilities, ignoring the rest.
- This helps in generating more focused responses while still allowing for some diversity.

Top-P

- Also known as nucleus sampling.
- Considers a **cumulative probability** distribution and includes tokens until the cumulative probability exceeds a certain threshold p .
- Allows dynamic adjustment based on the changing probabilities of tokens.
- Assuming $p=0.8$, top-P sampling will consider tokens until the cumulative probability exceeds 0.8.

yaml

 Copy code

```
Token A: 0.3 (considered)
Token B: 0.2 (considered)
Token C: 0.15 (considered)
Token D: 0.1 (not considered, as cumulative probability exceeds 0.8)
```

Maximum Length

- Maximum length
 - Set the maximum number of tokens of the output.
 - With GPT-3.5, the maximum allowed is 2K (2048) or approximately 1500 words.
 - Token calculator:
<https://platform.openai.com/tokenizer>

GPT-3.5 & GPT-4 GPT-3 (Legacy)

We, the citizens of Singapore, pledge ourselves as one united people, regardless of race, language or religion, to build a democratic society based on justice and equality so as to achieve happiness, prosperity and progress for our nation.

Clear Show example

Tokens	Characters
45	239

We, the citizens of Singapore, pledge ourselves as one united people, regardless of race, language or religion, to build a democratic society based on justice and equality so as to achieve happiness, prosperity and progress for our nation.

TEXT TOKEN IDS

Stop Sequence

- Stop Sequences
 - Tell the model when to stop generating an output.
 - For example: If you want a one-sentence answer, you can use “.” as the Stop sequence.
 - Alternatively, for a one-paragraph answer, you can use New Line “\n” as the Stop sequence.
- Useful when trying to generate a dialogue, Q&A or any kind of structured format.
- Most of the time, it is not used frequently.

Frequency and Presence Penalties

- A lot is done under the hood to make sure that it does not just generate the same text again and again.
- One of the ways is by automatically penalizing tokens that is already used.
- API provides two penalty controls.
 - Frequency penalty
 - Penalizes tokens based on how many times the tokens appeared in the text i.e. more time the tokens appear, the more they are penalized
 - Presence penalty
 - Penalizes tokens based on whether the tokens have already appeared in the text.

Chunking & Document Splitting

- Splitting a long document into smaller chunks that can fit into the model's context window.
- Splitting documents into smaller chunks is important and tricky as we need to maintain meaningful relationships between the chunks.
- When you chunk data, overlapping a small amount of text between chunks can help preserve context. We recommend starting with an overlap of approximately 10%.
- For example, given a fixed chunk size of 256 tokens, you would begin testing with an overlap of 25 tokens. The actual amount of overlap varies depending on the type of data and the specific use case, but we have found that 10-15% works for many scenarios¹.

Reference: 1. <https://learn.microsoft.com/en-us/azure/search/vector-search-how-to-chunk-documents>



Chaining

Prompt Chaining

- One of the important prompt engineering techniques is to break tasks into its subtasks.
- The LLM is prompted with a subtask and then its response is used as input to another prompt, creating a sequence of prompts that lead to the final results.
- Prompt chaining is like assembling a series of building blocks to construct a complete solution.
- Instead of overwhelming the LLM instance with a single detailed prompt, we can use the sub-prompts to guide it through multiple steps, making the process more efficient and transparent.

Prompt Chaining Advantages

- Simplified instructions
 - Write less complicated instruction instead of trying to express a complex task in a single prompt.
 - Break it down into smaller, more straightforward steps.
 - Increase the chances of getting accurate results.
- Help troubleshooting
 - Pinpoint specific prompt in the chain that needs adjustment when LLM has difficulty getting accurate or desired result.
- Incremental validation
 - Check LLM's output in stages rather than wait till the end.
 - Allow to assess the correctness of responses as it progress through the prompt stages.


Chain Types

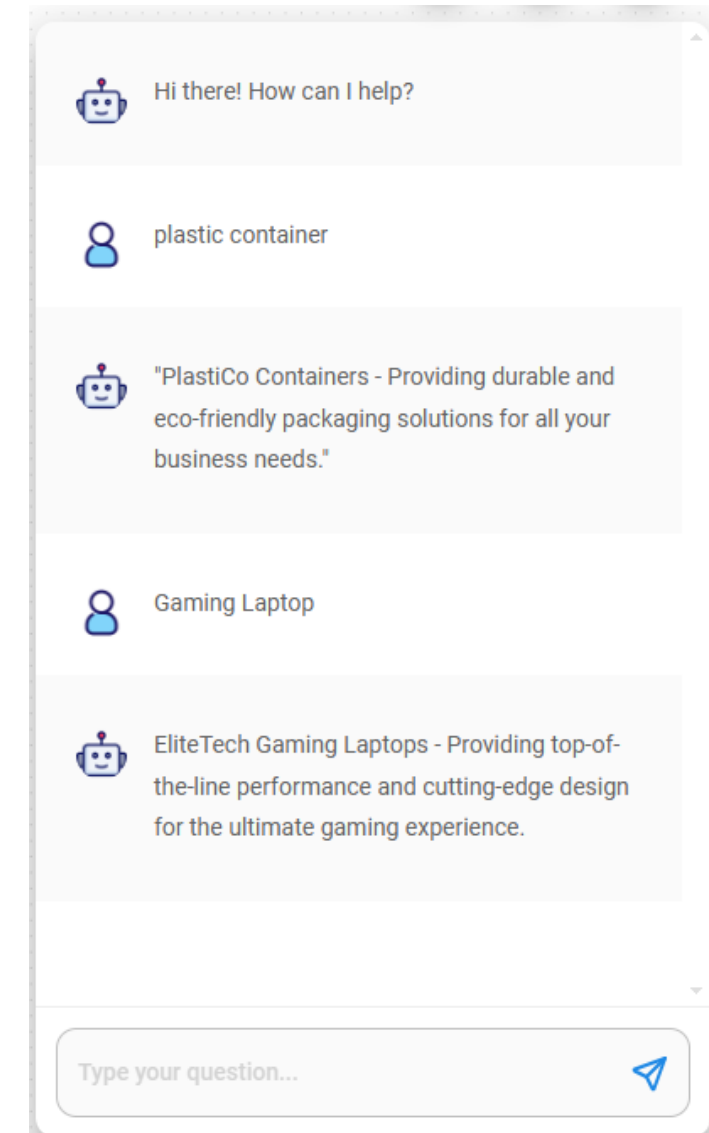
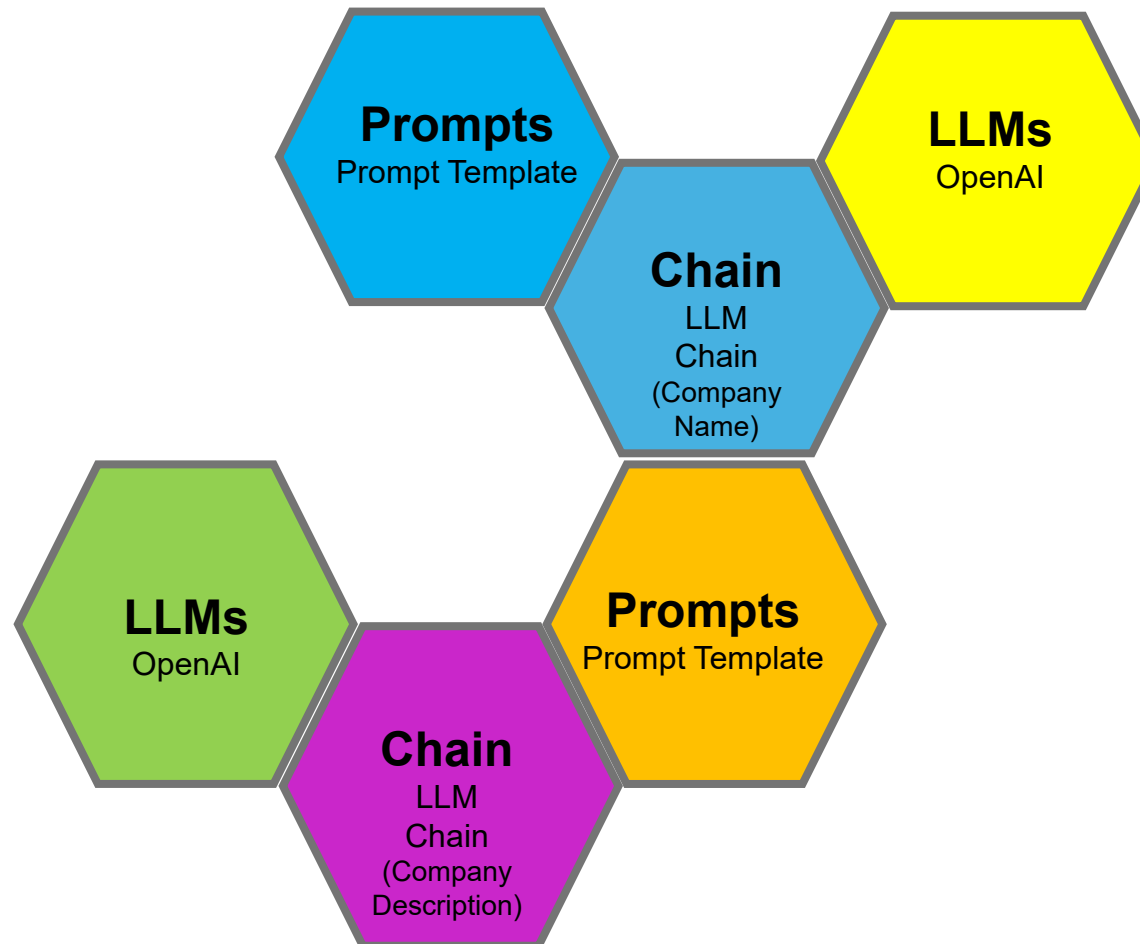
- Simplest Chain
 - User Input → Prompt Template → LLM Model
- Sequential Chain
 - Simple sequential chain: handle a single input and output.
 - Sequential chain: manage multiple inputs and outputs simultaneously.
- Router Chain
 - Allows routing inputs to different destination chains based on the input text.
 - Allows the building of chatbots to handle diverse inputs.



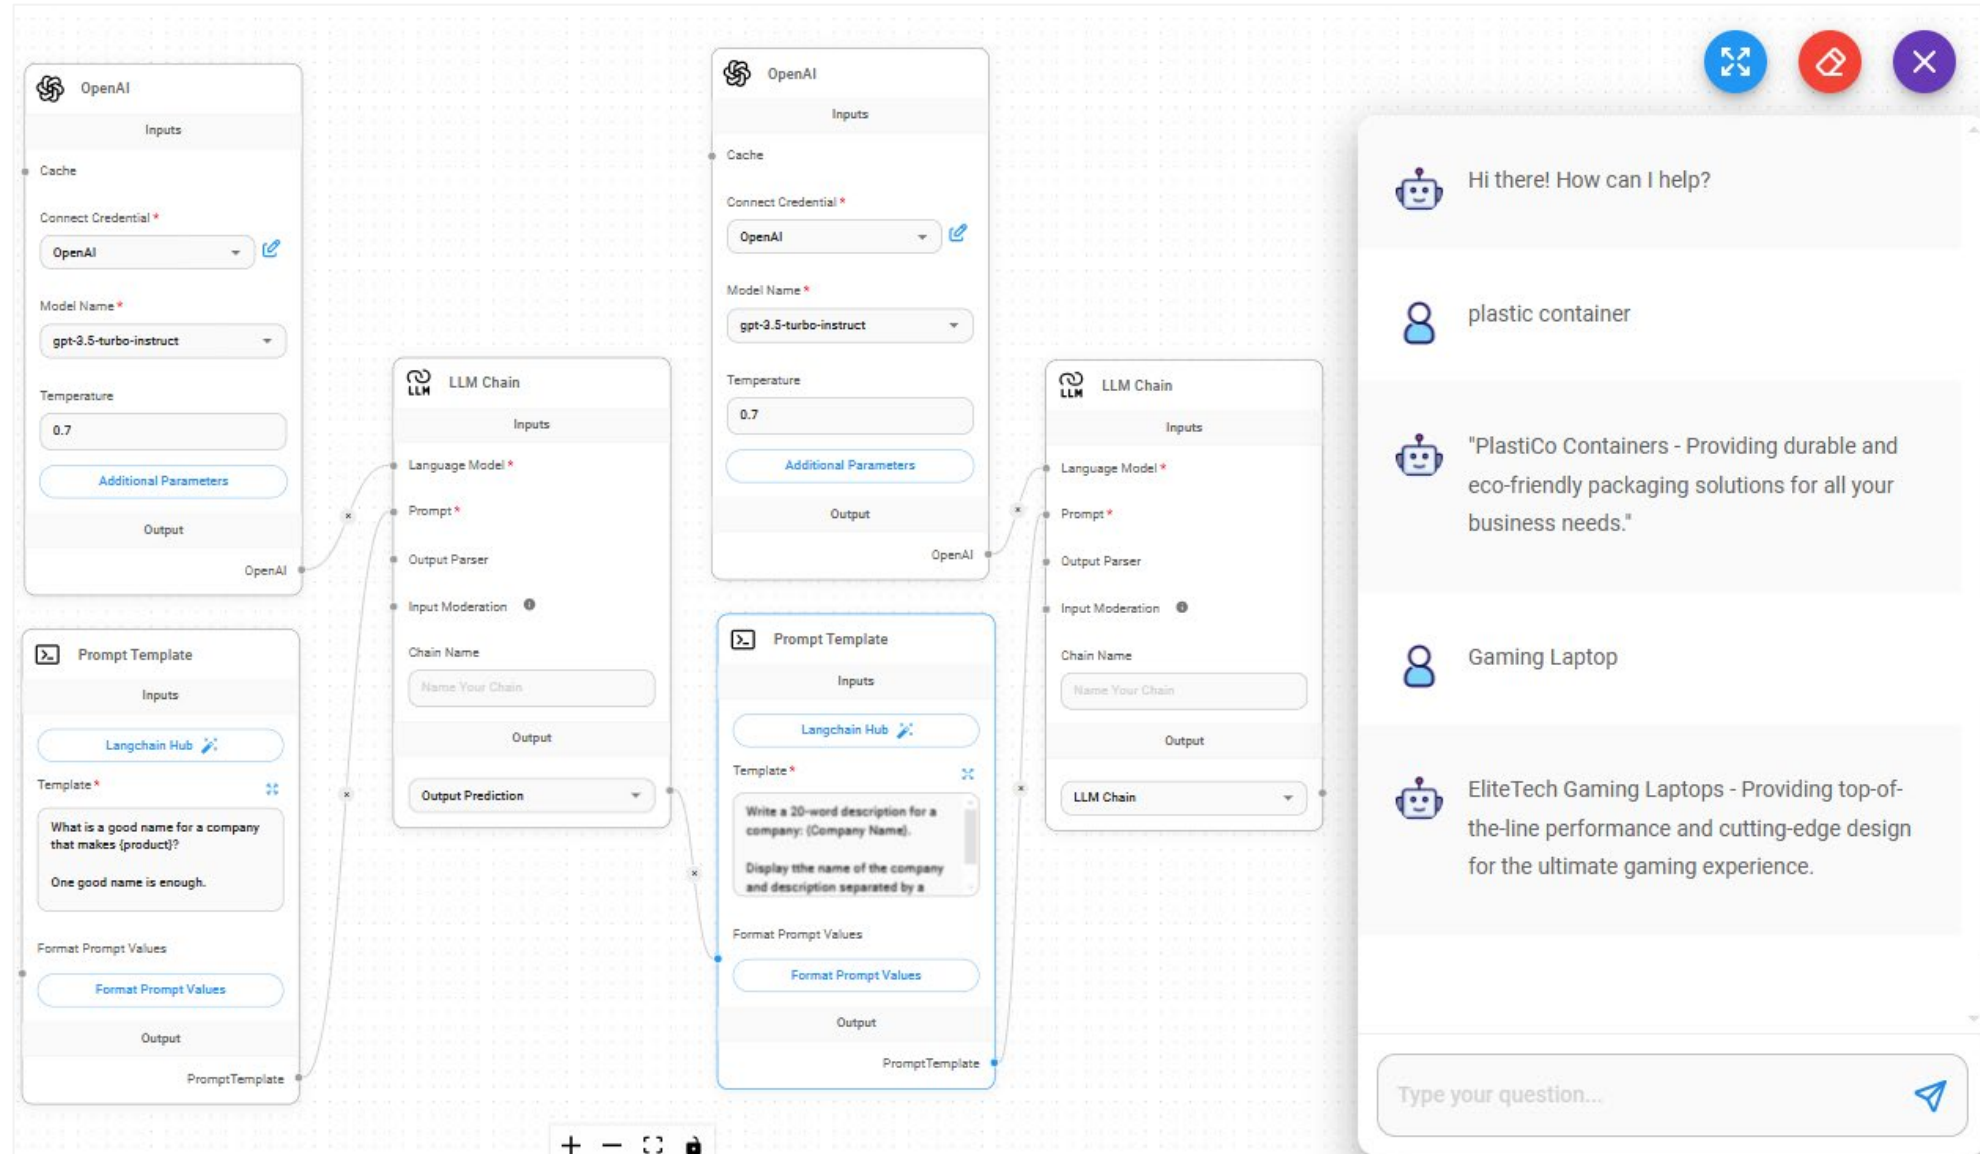
Activity: Sequential Chain

Activity Map (Sequential Chain)

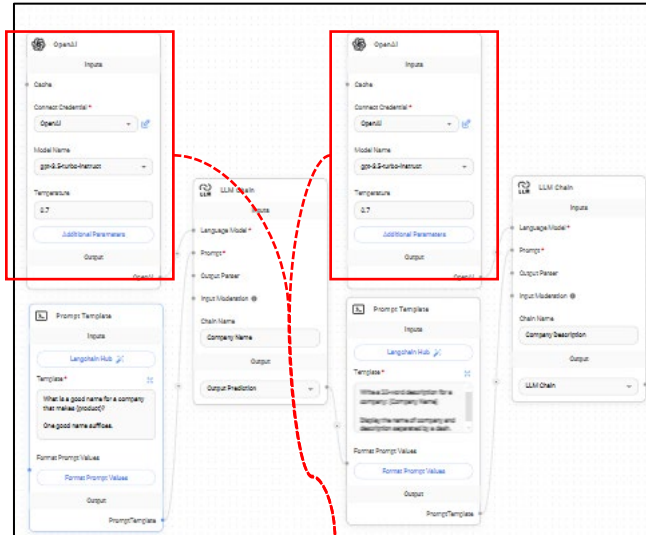
- You will need the following:
 - ☒ OpenAI API Key [ OpenAI]




Sequential Chain



Sequential Chain




OpenAI

Inputs

Cache

Connect Credential *

OpenAI

Model Name

gpt-3.5-turbo-instruct

Temperature

0.7

Additional Parameters

Output

OpenAI

Max Tokens

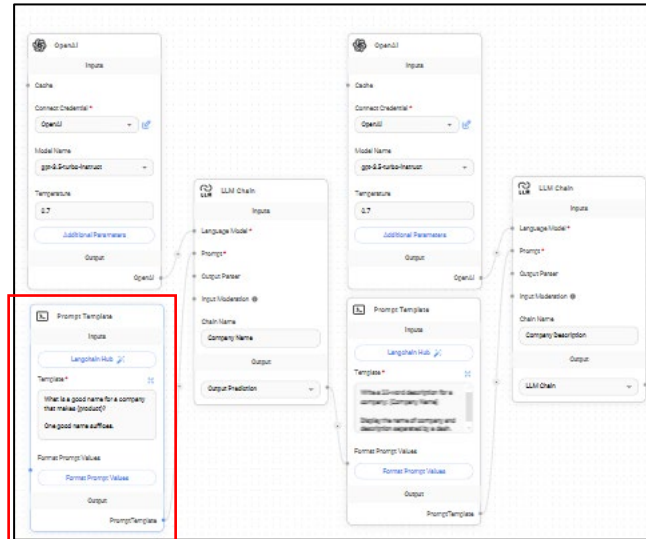
Top Probability


Best Of

Frequency Penalty

Presence Penalty

Sequential Chain




Prompt Template

Inputs

Langchain Hub

Template *

What is a good name for a company that makes {product}?

One good name is enough.

Format Prompt Values

Format Prompt Values

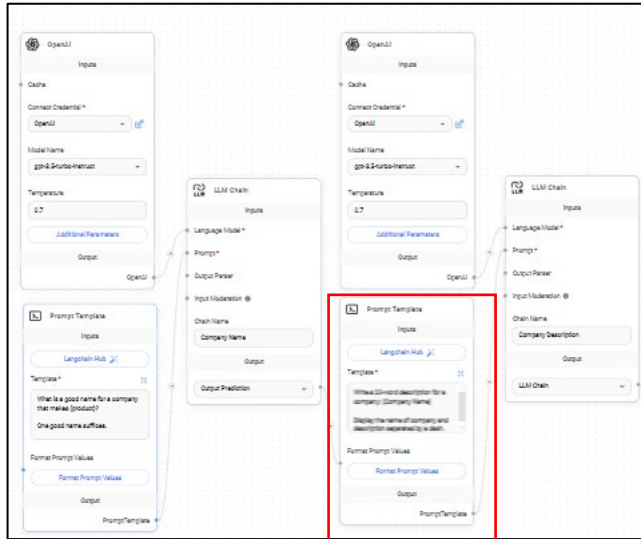
Output

PromptTemplate

Format Prompt Values

```
{ 1 item
  product : "{{question}}"
}
```

Sequential Chain



Prompt Template

Inputs

Langchain Hub

Template *

Write a 20-word description for a company: {Company Name}.

Display the name of the company and description separated by a

Format Prompt Values

Format Prompt Values

Output

PromptTemplate

Select Variable

question

User's question from chatbox

chat_history

Past conversation history between user and AI

file_attachment

Files uploaded from the chat when Full File Upload is enabled on the Configuration

LLM

Output Prediction from LLM Chain

Format Prompt Values

{ 1 item


Company Name : "{{llmChain_1.data.instance}}"

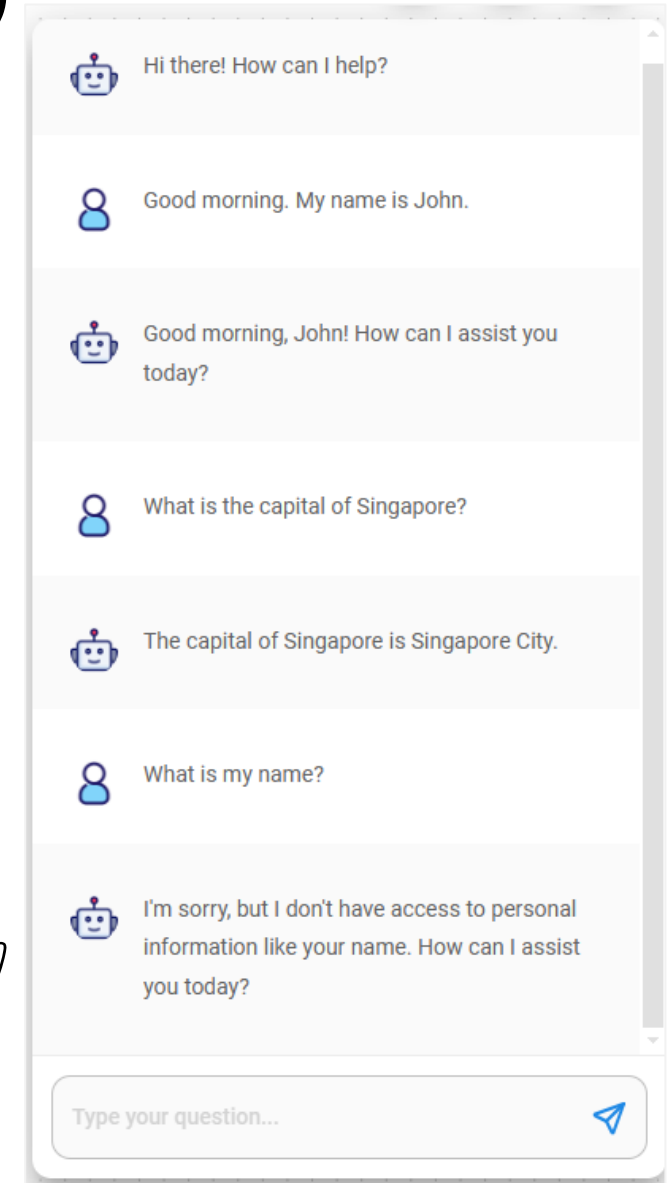
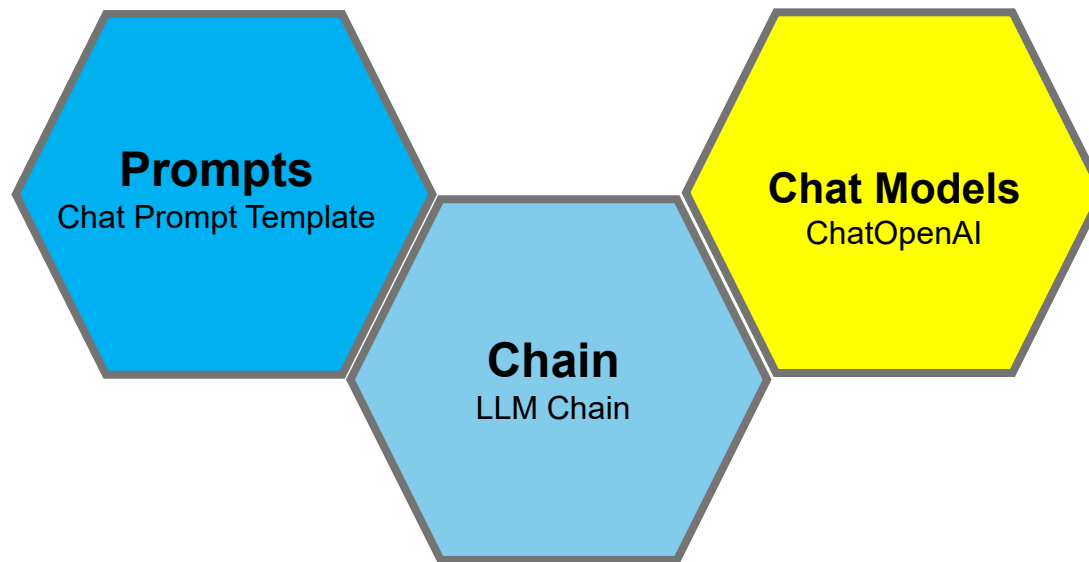
Reminder: Save your work before testing



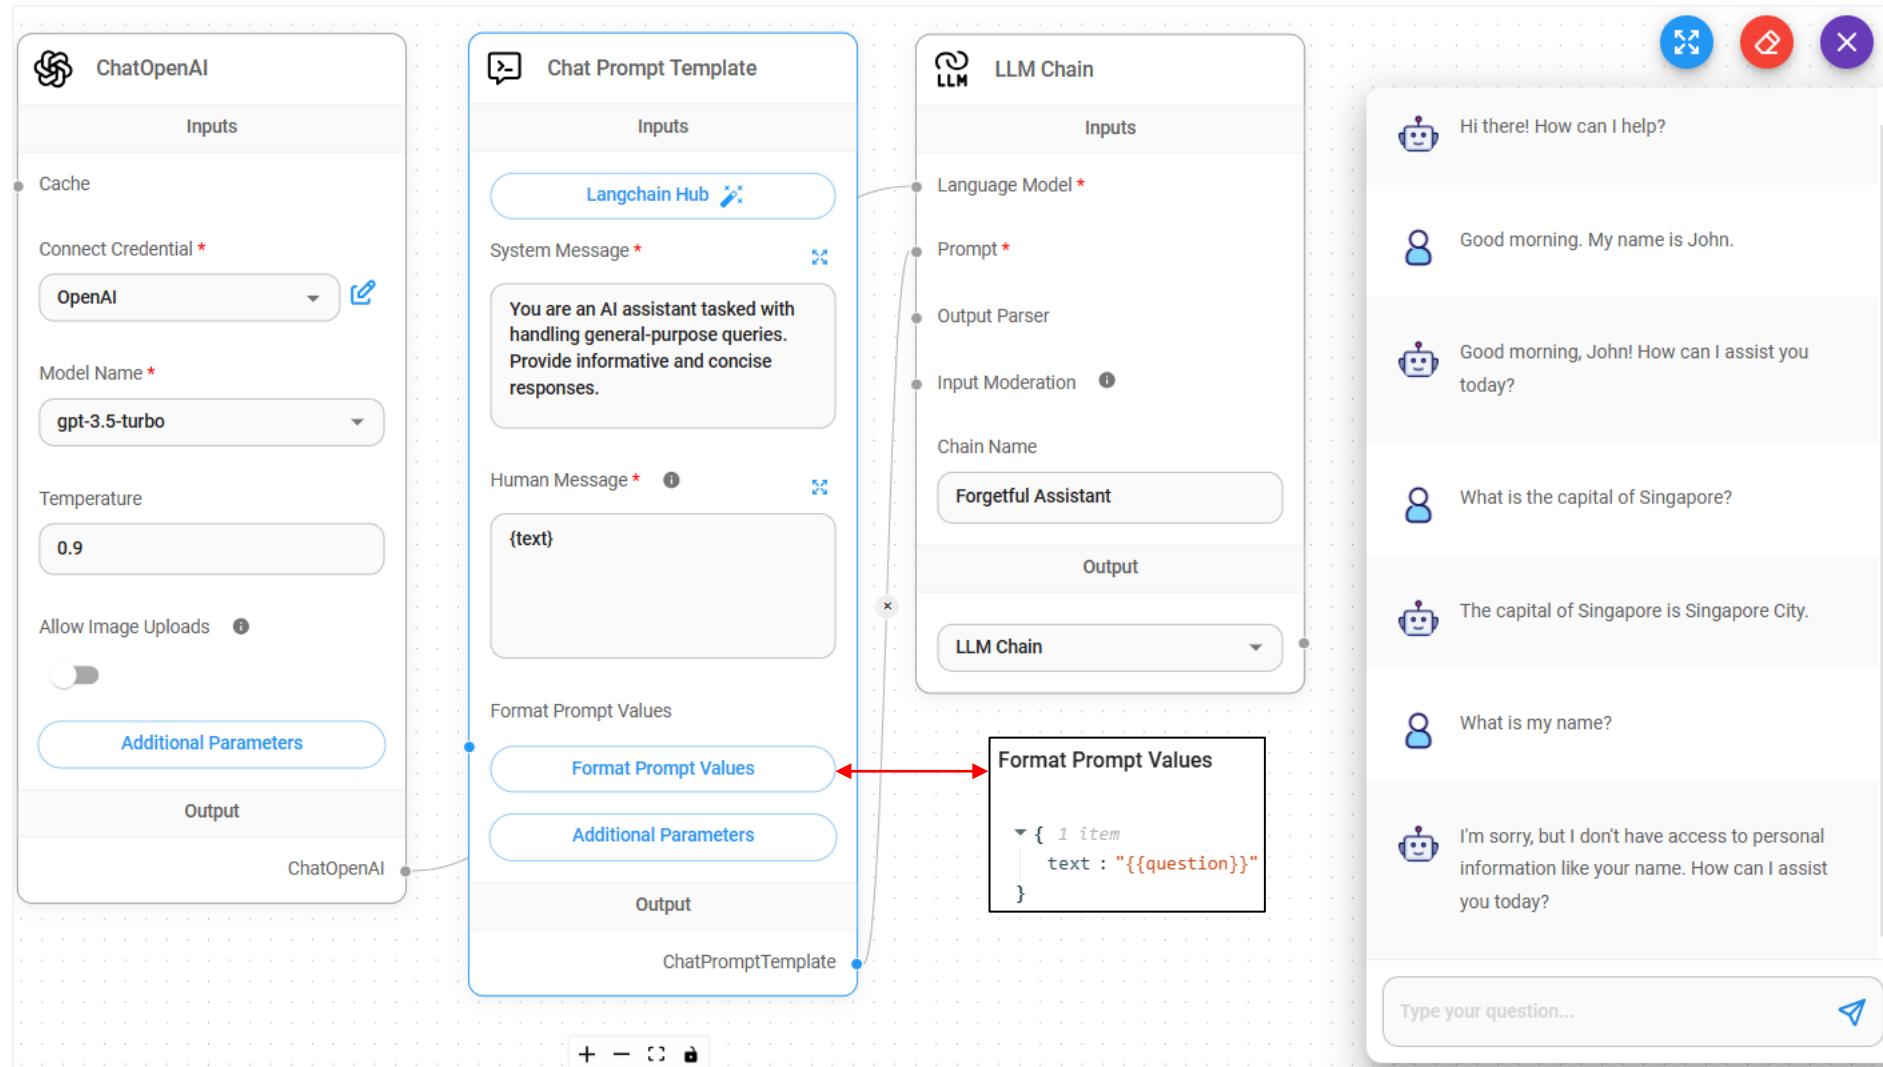
Activity: Memory

Activity Map (Forgetful Assistant)


- You will need the following:
 - ☒ OpenAI API Key [ OpenAI]

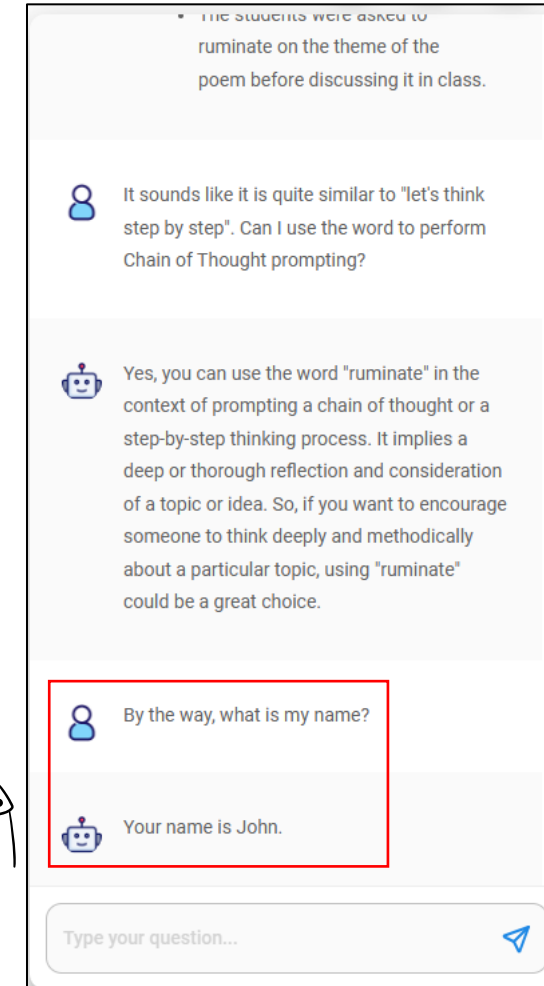
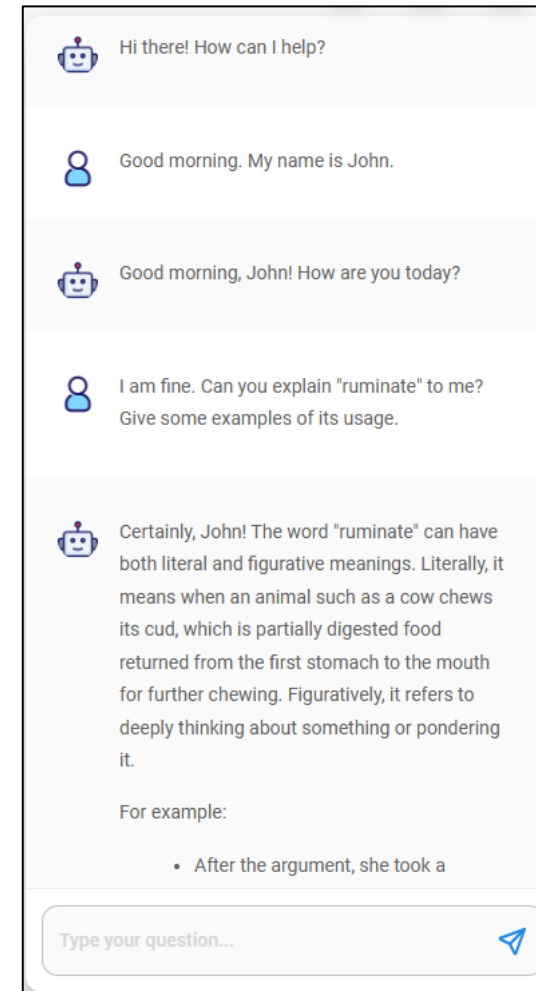



Forgetful Assistant

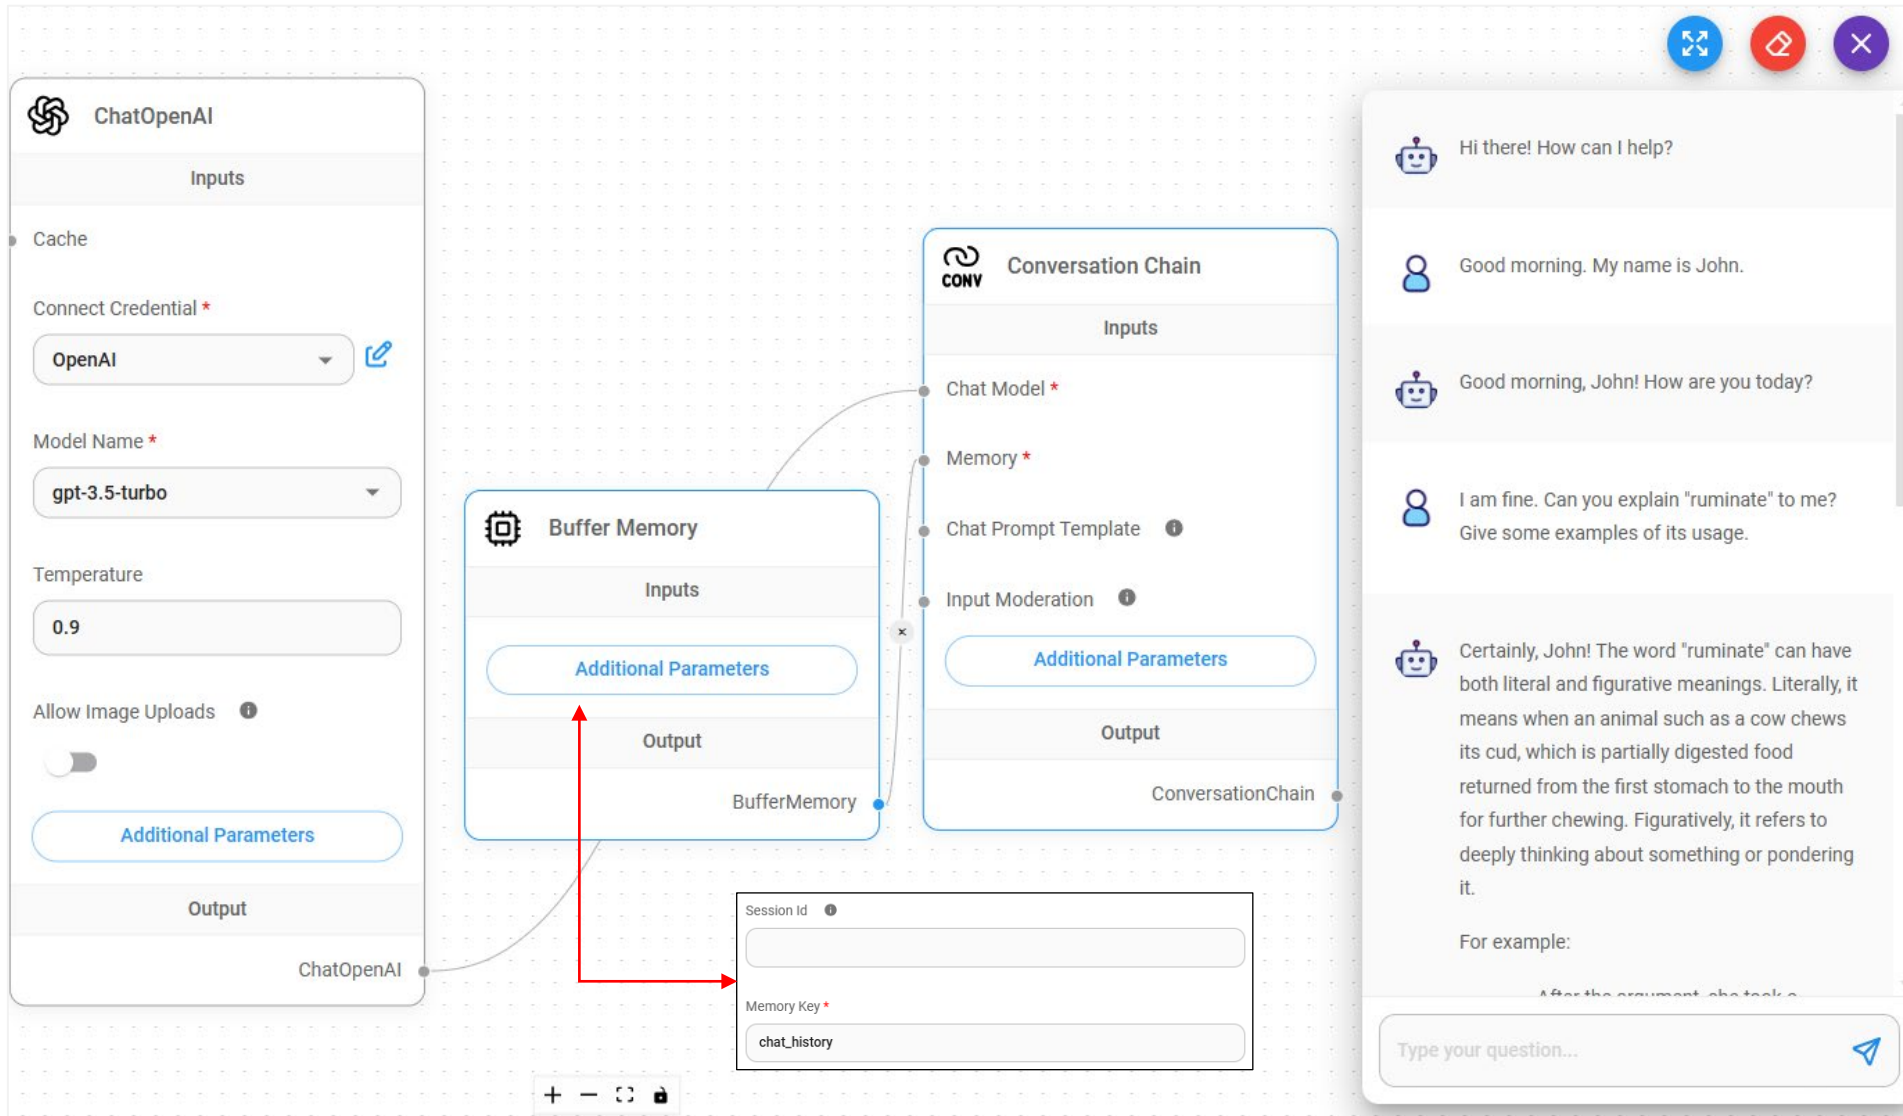


Activity Map (Attentive Assistant)

- You will need the following:
 - ☒ OpenAI API Key [ OpenAI]




Attentive Assistant

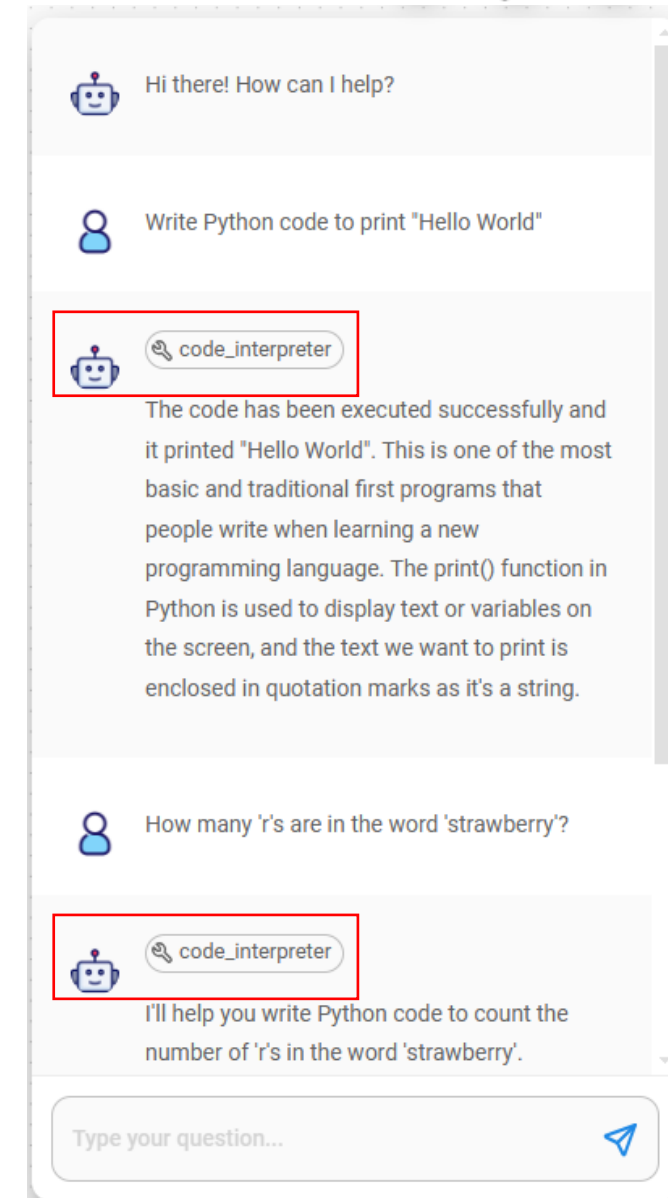
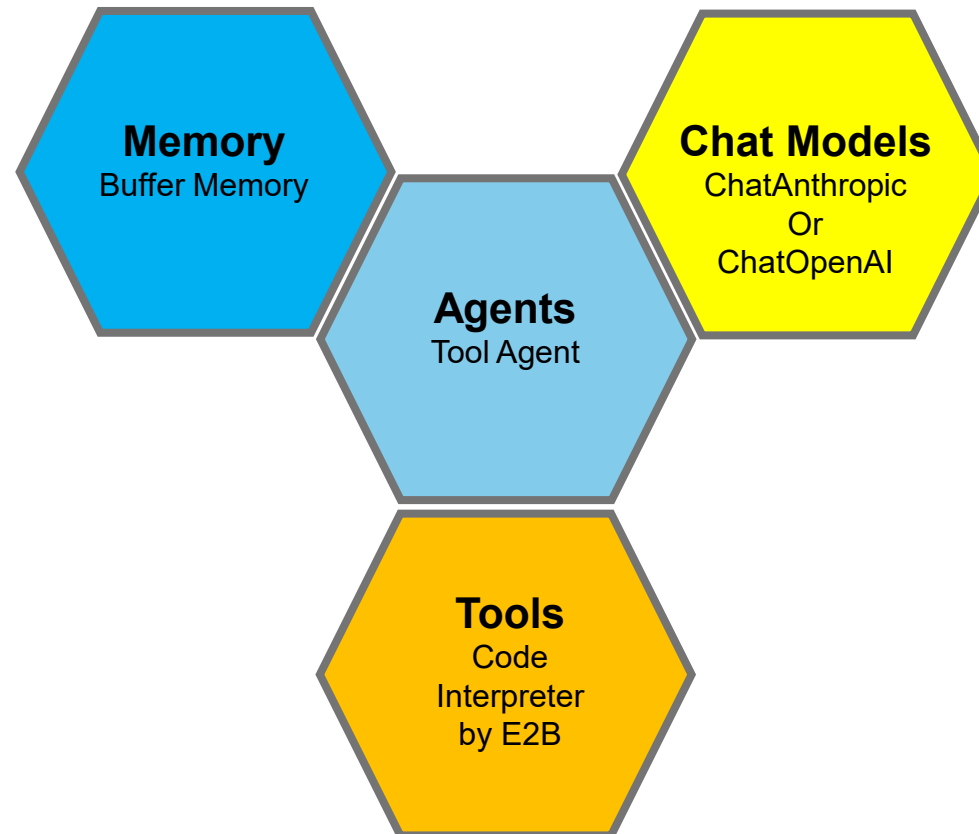




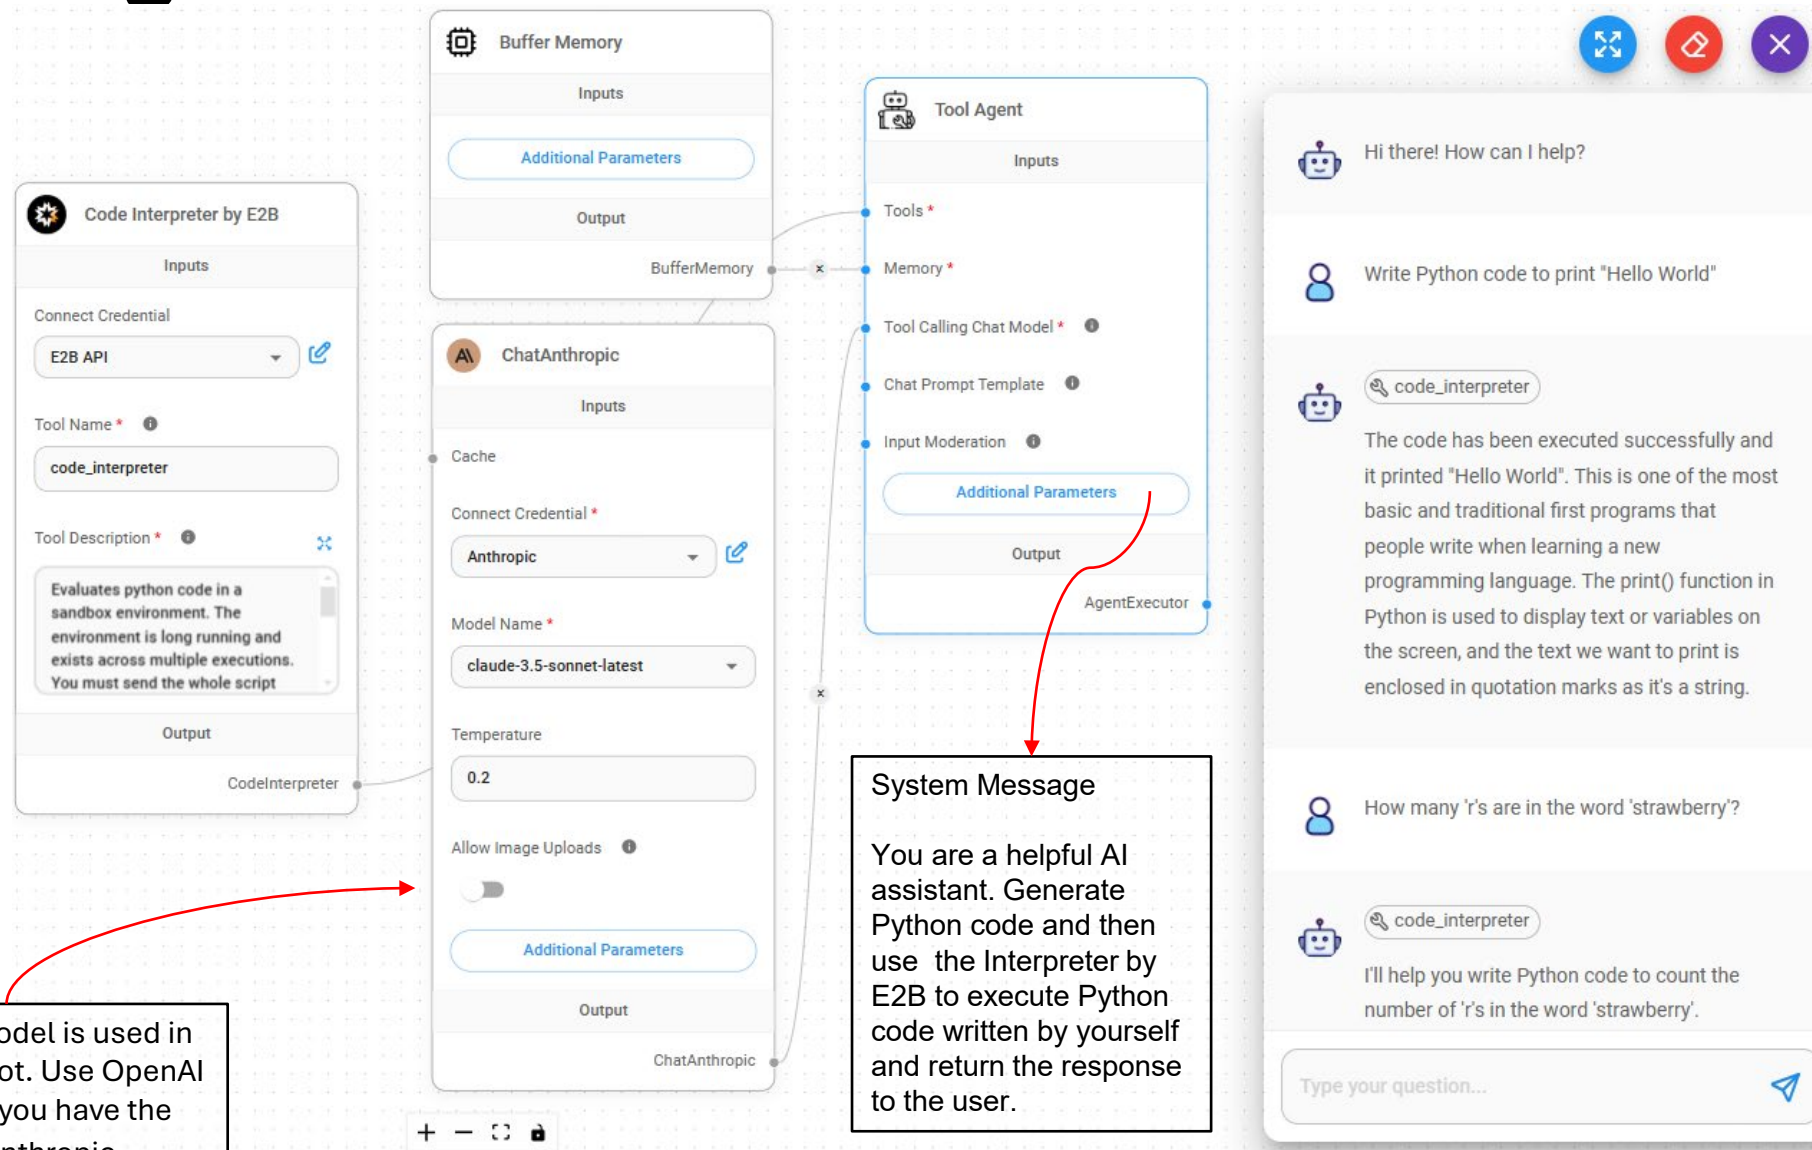
Activity: Agents

Activity Map (Tool Agent)

- You will need the following:
 - ☒ OpenAI API Key [ OpenAI]



Tool Agent



* Anthropic model is used in this screen shot. Use OpenAI model unless you have the API key from Anthropic.

Tool Agent Observations

1

Hi there! How can I help?

Write Python code to print "Hello World"

`code_interpreter`

The code has been executed successfully and it printed "Hello World". This is one of the most basic and traditional first programs that people write when learning a new programming language. The print() function in Python is used to display text or variables on the screen, and the text we want to print is enclosed in quotation marks as it's a string.

Used Tools

```
{
  3 items
  tool : "code_interpreter"
  toolInput : { 1 item
    input : "print('Hello World')"
  }
  toolOutput : "Hello World "
}
```

2

How many 'r's are in the word 'strawberry'?

`code_interpreter`

The code counts the occurrences of the letter 'r' in the word 'strawberry' using the string method count(). As we can see, there are 3 'r's in the word 'strawberry'. The letters are located at:

1. First 'r' in "str"
2. Second 'r' in "ber"
3. Third 'r' in "rry"

Type your question...

Used Tools

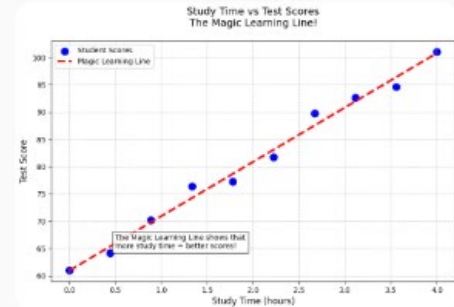
```
{
  3 items
  tool : "code_interpreter"
  toolInput : { 1 item
    input :
      "word = 'strawberry' count = word.count('r') print(f'The letter 'r' appears {count} times in the word '{word}'.)"
  }
  toolOutput :
    "The letter 'r' appears 3 times in the word 'strawberry'."
}
```

3

Hi there! How can I help?

Draw a chart to explain linear regressions to a 10 years old.

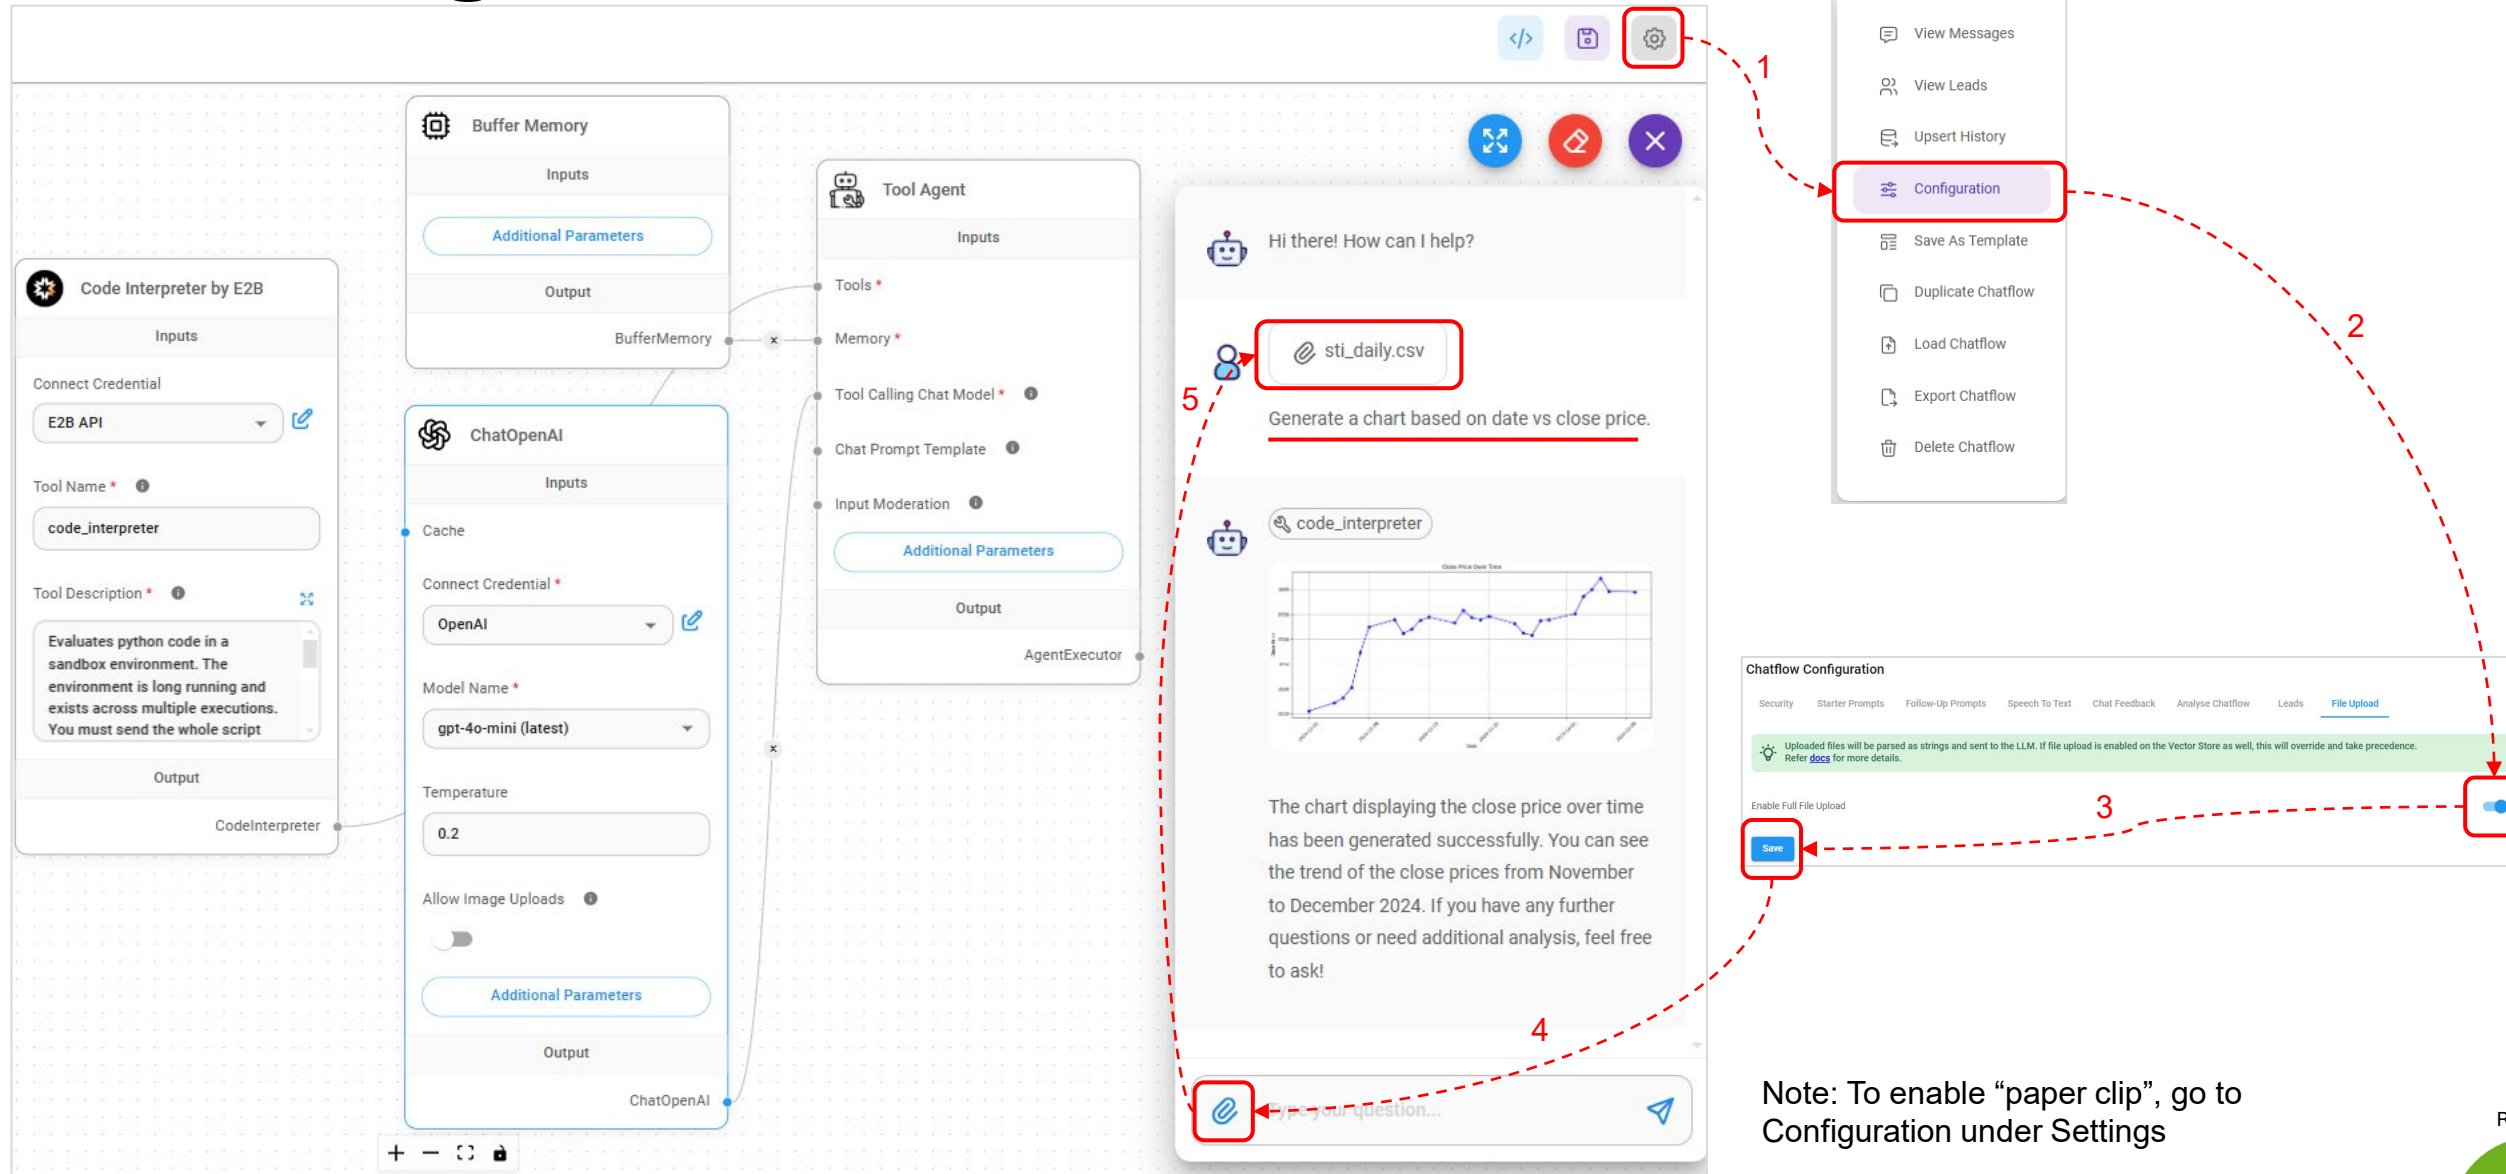
`code_interpreter`



I'll help create a simple and colorful chart to explain linear regression to a 10-year-old using an example they can relate to - the relationship between study time and test scores. I'll create some sample data and plot it with a clear regression line.

Type your question...

Tool Agent Observation



1 Settings icon

2 Configuration menu

3 Enable Full File Upload toggle

4 Save button

5 Paper clip icon for file upload

sti_daily.csv

Generate a chart based on date vs close price.

code_interpreter

The chart displaying the close price over time has been generated successfully. You can see the trend of the close prices from November to December 2024. If you have any further questions or need additional analysis, feel free to ask!

Chatflow Configuration

Security Starter Prompts Follow-Up Prompts Speech To Text Chat Feedback Analyse Chatflow Leads **File Upload**

Uploaded files will be parsed as strings and sent to the LLM. If file upload is enabled on the Vector Store as well, this will override and take precedence. Refer [docs](#) for more details.

Enable Full File Upload

Save

Note: To enable "paper clip", go to Configuration under Settings

Reference

- Flowise

<https://docs.flowiseai.com/>

Thank you!